

Sergei Ponomarev

Your Company's AI Through the Customer's Eyes

Five requests from your most demanding customer:

1. First, ask me what I need.
2. Then prove that your AI benefits me.
3. Next, let's agree on the rules.
4. Talk to me like a human being.
5. Check yourselves through my eyes.

aibusiness.vc

2026

Contents

Introduction. I Am Your Most Demanding Customer	3
Chapter 1. Ask the Customer First.....	6
1.1. My Needs. There Are More of Them Than You Think	6
1.2. Levels of Collaboration. From "Respondent" to "Co-author"	8
1.3. How to Ask So I Don't Lie.....	10
Chapter 2. Why Would Your Customer Need Your AI?	13
2.1. My Three-Number Test	13
2.2. The AI Guillotine: A Presumption of Corporate Inefficiency	14
2.3. Time Is Relative: Your Instant and My Eternity	16
Chapter 3. Agreeing on the Rules.....	18
3.1. AI Policy: How Things Are Done Here.....	18
3.2. The AI Service Passport: The Standard Your AI Works By	19
3.3. The AI Receipt: What Actually Happened	21
3.4. Let Me Show You: Policy, Passport, Receipt	23
Chapter 4. Talk to Me Like a Human Being	27
4.1. Speak Plainly.....	27
4.2. AI as an Occasion to Care About Me	28
4.3. People, Not "Personnel Costs"	30
Chapter 5. How to Evaluate AI	33
5.1. Build an Evaluation System, Not a Collection of Likes.....	33
5.2. The Mystery Shopper: Hire Me	34
5.3. Living Rules: Revise Them Together With Me	35
Conclusion. Free Up My Chair	38
Appendix. The Demanding Customer's Final Checklist	39
About the Author	42

In Word: press Ctrl+A, then F9 to update the table of contents.

Introduction. I Am Your Most Demanding Customer

Hello. I am your worst nightmare: your most demanding customer. Pleased to meet you.

I am the person who pays you money, the one your company works for.

Lately, though, I have noticed that many businesses are adopting AI without asking my opinion at all.

They are busy choosing an AI model, hiring a vendor, calculating the cost, estimating how much they can save on staff, and so on. Fair questions, no argument. But the list almost never includes the question they should start with: what will I, your customer, think about all this?

However advanced, modern, economical and marvelous your AI may be, if I, a mere meatbag, do not like it, you will have problems. According to MIT, 95% of generative AI pilot projects end with zero results. The reasons for failure vary. This book is about one of the systematically underestimated ones: companies design AI around the technology, not around the customer.

Jeff Bezos understood this better than anyone. He built Amazon around customer obsession and kept one empty chair at the table in important meetings. In that chair sat "the demanding customer," the most important employee (!!!) in the room. Decisions were always made with an eye on that chair.

You are in luck. If you are reading this guide, it means I have come to you, bringing the customer's view of AI adoption in your company. I will help you look at the task through your customers' eyes. Use AI as a tool and an occasion to show you care about them, and they will repay you in hard cash.

And to make it easier for you, I will take this heavy role upon myself: I will be your collective, universal and, fair warning, maximally demanding customer. If I like it, everyone will. I am going to tell you unpleasant things, for which I apologize in advance! I will explain what I need, what irritates me, what I am willing to pay for, and what will make me leave and never come back. You can, of course, ignore me. That is your right. Just don't act surprised later.

What gives me the right to think I can pull off this role? My rather unusual professional experience at the intersection of business, academic research and civic policy.

My name is Sergei Ponomarev, and I hold a PhD in political science. For more than 20 years I have been studying how digital services are created and delivered in practice. In 2014 I defended a dissertation on e-government adoption, completed a fellowship on government openness at the Kennan Institute in Washington, DC, and taught at a university for 14 years.

For 7 years I led applied research at a think tank: I ran a nationwide monitoring program for government and municipal services, conducted independent quality assessments and mystery shopping audits to uncover the real problems people face. I developed an Information Openness Standard for public authorities. I worked with several human rights organizations and projects focused on consumer rights.

For 10 years I have been running trainings and seminars, teaching and advising entrepreneurs on service quality.

For the past year and a half I have been building AI agents hands-on and running the portal aibusiness.vc.

In other words, I spent two decades making sure that the services people receive are open, understandable and convenient. First in the public sector, now in business.

Why this matters right now. A large part of what I am going to ask of you in this guide will soon stop being a polite request.

On August 2, 2026, the transparency requirements of the EU AI Act¹ start to apply. It is important not to mix up the different obligations here. Providers of systems that interact directly with people must ensure that from the start of the first contact the person understands they are talking to AI, unless that is already obvious. Providers of generative systems ensure machine-readable labeling of the synthetic content covered by the law. Companies deploying AI separately disclose the use of emotion recognition and biometric categorization, and explicitly label deepfakes and certain texts on matters of public interest published without proper human oversight. In other words, the law does not demand that everything AI helped create be labeled the same way: the specific obligation depends on the company's role and the scenario. Fines for violating the transparency requirements run up to 15 million euros or 3% of turnover. In the US, state laws on specific AI scenarios are also in force, but their scope varies and is usually narrower than the European one.² So from here on I will distinguish three things: what the law requires, what I propose as a voluntary consumer standard, and what belongs to technical implementation. Not everything in this brochure is required by law, but much of it is worth doing before your competitors do, and with a smile.

Many entrepreneurs wonder how corporate AI should properly be governed: after all, it is a completely new phenomenon in history. The principle is very simple. Wherever and whenever AI interacts with me as a customer, a perfectly traditional and well-known consumer framework kicks in. It rests on simple, straightforward rules, even if the service is now delivered not by an employee but by a wonder machine.

To use a service, I first have to learn that it exists. So you need an AI policy: how and what kind of AI you use to work with me, and what I am guaranteed along the way.

When I receive a service, I want to receive it the way it was promised. So every AI service needs a standard, or passport: what exactly a specific bot can do and promises, so that I have something to compare its actual output against.

And at the end of the interaction, hand me a receipt. The entire consumer movement started a hundred years ago with the demand to "always take the receipt," because the receipt is the only key to an appeal if something goes wrong. By the way, the receipt protects you too, so do not neglect it: otherwise our dispute turns into "the customer's word against tech support's word."

So everything you and I agree on in this brochure will stand on three pillars: the company's AI policy, the service passport and the receipt. None of them works without the others. Without the policy, the passport has nothing to lean on. Without the passport, the receipt has nothing to be checked against. And the receipt protects us both.

We will take all of this apart in detail across five chapters, which I have framed as five of my requests:

- First, ask me what I need.
- Then prove that your AI benefits me.
- Next, let's agree on the rules.
- Talk to me like a human being.
- Check yourselves through my eyes.

And at the end I will hand you a checklist of all my requirements for your AI.

An important clarification. I am offering a consumer framework for designing AI services, not a legal opinion. Which provisions are mandatory depends on the country, the industry, the data, the company's role and the risk level; the legal and technical fields should be filled in by the relevant specialists.

Read with a pencil in hand, argue with me, adapt things to your company. I am demanding but fair: every request will be explained and backed by a method for meeting it.

Portal: aibusiness.vc

For collaboration: info@aibusiness.vc

Let's begin. And we begin, naturally, with me.

¹ Regulation (EU) 2024/1689 (EU AI Act), Article 50: eur-lex.europa.eu/eli/reg/2024/1689. European Commission guidance on Article 50 of 20.07.2026: digital-strategy.ec.europa.eu/en/faqs/transparency-obligations-under-article-50-ai-act

² Examples of US state laws with varying scope: California SB 243 (2025), companion chatbots only, not ordinary customer-facing bots; Utah Code Title 13, Chapter 77, disclosure in certain generative AI interactions and a limited safe harbor; Maine Title 10, §1500-DD, disclosure where a commercial chatbot could be mistaken for a human; Colorado SB 26-189 (2026), automated systems in consequential decisions, core requirements from 01.01.2027. The applicability of each law has to be checked separately.

Chapter 1. Ask the Customer First

You would not believe how rarely this happens.

Companies spend millions on AI adoption and almost never do the cheapest and most useful thing: they do not ask the customer what they actually need.

And it is not as if you had to go far to find us. We call you, write to you, browse your websites, leave reviews. We are practically shouting about our needs every single day. You just have to know how to listen.

In this chapter I will explain what customer needs are made of and how best to work with us at the AI adoption stage. And yes, how to ask so that I do not lie to you. I lie quite often, out of politeness. We will talk about that too.

1.1. My Needs. There Are More of Them Than You Think

Let me start with an uncomfortable truth about our relationship. When I come to you for a service, we are not equals. You made the rules. You are more experienced than I am, because you sell hundreds of services every day. You designed the interface. You decided what your bot can do. Nobody ever asked me whether I want to play by your rules. Sound familiar?

So yes, we look at your business differently. You judge a service by the result: order delivered, case closed, box ticked. I judge it by the process. If the order was completed but I spent half an hour begging your bot to transfer me to a live agent, I consider the deal a failure and your company terrible.

To avoid that, memorize my basic needs. They fall into three large groups. This classification is not from a textbook: it comes from years of service quality research involving thousands of real customers. People everywhere want the same things. Only the scenery and the operators change.

First group: functional needs. Where, when and how.

The place of service. Often the place of service is the company's office. Everyone has a rough idea of what a good office looks like. You do, right? I would genuinely love to hear it.

With AI, the assistant also needs its own "place," which is usually a channel and an interface. Where does your assistant live: on the website, in the app, in a messenger? I am comfortable where I already am, not where it was cheapest for you to put it. What does the AI assistant's "office" look like: is the font large, is the button visible, how many menus do I have to click through to ask a simple question? A cluttered chat with three carousels and ten buttons is like a store where the goods are piled in a heap and the salesperson hides in the back room.

Time. When the service is available and how long it takes. Here AI has a natural trump card: it does not take lunch breaks, does not get sick and does not go on vacation. Can I solve my problem at night on a Sunday? Wonderful. But opening hours are not the whole story.

There is also the duration of the service: how long the process itself takes. A bot that replies instantly but runs me in circles for forty minutes is available around the clock. Around the clock, and badly.

The manner of delivery. How many actions I have to perform. Every extra button, every "please register first," every "fill out this 12-field form" filters out people like me. We are not lazy. We simply remember that your competitor needed three clicks.

You have surely heard of the "one-stop shop" principle. As a customer, I want to solve all my problems in one place. If the bot takes my order but cannot answer my questions, cannot track the order status, and cannot accept a complaint because I have to "go" to a separate form on the website, then you do not have a smart assistant. You have a convoluted quest game.

And yes, staff: a live human must remain in the loop. Not hidden behind seven buttons, but reachable in one tap. I will have a separate conversation with you about this in Chapter 3.

Second group: informational needs. What is going on?

Accessibility. Before I use your smart bot, I first need to understand what it can do, what it cannot do, and what authority it has: can it process a refund or only chat about one? How long do I wait for a human if the bot fails? An accessible service means simplicity and clarity for the customer, without fog. Uncertainty irritates me more than rejection: a refusal I can understand, fog I will never forgive. Remember this need of mine: in Chapter 3 it will grow into an entire document, the service passport.

Confidentiality. What happens to what I tell you? Where does the chat log go, who reads it, does your AI train on it, who do you share it with? As long as you stay silent, I assume the worst. That is the kind of person I am, and the law, by the way, is on my side. The answers to these questions of mine should not live in your lawyer's head. Their place is in the company's AI policy, which we will also get to.

Reliability and safety. Can I trust what your AI says? If the bot quoted one price and a different one appeared at checkout, that is not a "neural network hallucination." That is misleading the customer. Which is exactly why I need a receipt from your AI agent. Your AI's promise equals your company's promise. Not ready to answer for its words? Then how can it represent you?

Contact and feedback. Can I reach you, do my problems get solved, what happens to complaints? Note: handling problems and complaints is my informational need, not your "negativity management." I need information: was I heard, what is being done, when will it be fixed. Silence after a complaint infuriates me more than the problem itself.

Third group: emotional needs. Do not drive me crazy!

This is where all those synthetic bots fall apart. I do not demand smiles from AI, but politeness, empathy and understanding are mandatory.

Respect for dignity. Human dignity is hard to describe in words, yet it registers instantly in your gut the moment it is violated. If the bot has not understood the question on the third try, it should offer to call a human, not keep cheerfully answering off-topic. If the bot failed, it is better to apologize like a person, not to "regret any inconvenience caused."

Precision and accuracy. A precise service means getting it right the first time. Not after I have repeated myself twenty times, but the first time. In line with the standard, or the AI service passport. Do not mangle my name, my address, my grandmother's maiden name or my dog's name. Do not lose my order. Do not answer a question I never asked. Accuracy is a form of respect, and I register its absence instantly.

Satisfaction. The final feeling from the interaction: my problem is solved, not "the ticket is closed." You would be surprised how many companies close tickets without solving the problem. The bot has reported success, the statistics look great, and I am left alone with my issue. Know what I do next? I write a review. That review.

Now, lest I be accused of consumer terrorism, an important disclaimer that will save you money. Not all of my preferences need to be researched.

Some of them have long been written into law: the right to full information about a service, to safety, to protection of my economic interests, to a working complaints procedure, the right to a receipt. That is the baseline, enshrined in consumer protection laws. You do not need focus groups to discover it. You just need to comply with it.

Everything above the baseline is tied to your specific service, your audience and your market. That is what you should ask me about.

Write down this life hack. My preferences are perfectly visible in my behavior, in the traces your AI already records: at which step I abandon the dialog, which question I ask ten times, at what hour I storm into the chat, and after which answer I leave without a word. It is all already in your logs. You just have to look.

And one last thing. Write it somewhere in large letters. I am not interested in your internal regulations. I do not care which model you use, what your pipeline looks like, how your CRM is set up or what KPIs your support team has.

I judge only what is on my side of the counter: my functional, informational and emotional needs. Judge yourselves by them too. That is just how I am. And by the way, so are you, whenever you buy from someone else. Evaluation through the customer's eyes always follows the external procedures, not the internal ones. Your kitchen is your business.

How exactly to watch, ask and involve me, I will explain in the next two sections.

1.2. Levels of Collaboration. From "Respondent" to "Co-author"

Let me reassure you right away: I do not expect you to dance around me over every button. I am an adult, and I understand you run a business, not a hobby club.

So here is the deal. The depth of our collaboration depends on three things:

how much time you have,

how much you want me to like the result,

what is at stake.

A bot that tells me your opening hours is one thing. An AI that decides whether to approve my installment plan is quite another. The higher the stakes, the higher you climb the ladder. And the ladder is simple, just three steps.

Step one: ask me first.

A survey. A short survey of your customer base. Two or three questions in the chat after an interaction. A dozen phone calls. A week of time and a little money.

What you get in return: my priorities instead of your imaginary KPIs. Here is a typical story. A company plans to deploy a voice robot for taking orders, spends six months choosing a vendor, haggles over the price. Then someone finally thinks to ask the customers. It turns out we were basically fine with how orders were taken. What drove us mad was something else: nobody could tell us where the order was and when it would arrive. That is an entirely different project. And, incidentally, a cheaper one.

When this is enough: simple services, low stakes, routine questions. The bot tells people the schedule, order status, delivery terms? Step one, and do not overthink it.

Just, I beg you, ask before the launch, not after. A survey after the launch is not called "preference research." It is called a post-mortem.

Step two: let me try it.

A consumer test. This level is for higher stakes. Show me the prototype before launch. Sit five to seven customers in front of a screen, hand us your bot to tear apart, and stay quiet. Just watch where we stumble.

Let me tell you a secret about why this works better than internal testing. Your testers know the design, how things are supposed to work. They verify the plan, and everything checks out. I do not know how things are supposed to work. I type with typos, start from the middle, ask two questions in one message and drop in a voice note where nobody expected one. I have a talent for breaking other people's scenarios. We all have that talent, and there are millions of us.

What you get in return: a list of failures before the launch instead of after. Feel the difference in money. A failure found on a prototype costs you a prompt edit. A failure I find in the wild costs you a customer, a review and a social media pile-on.

When this is mandatory: the AI stands in the customer's main path. Ordering, booking, personal account, support. Everywhere I walk every day.

Step three: involve me.

Partnership. This level is for serious stakes: AI for the long haul, AI at the core of the business, AI whose decisions touch my money. One-off surveys are not enough here. You need a standing connection.

What it looks like. A panel of regular customers who see every new version first. A beta testers club with early access or a discount. Regular meetings between your team and living customers: we talk, you take notes. A customer council, if you want a formal name.

What you get in return, and pay attention here, because this is the second most valuable benefit in this entire guide. A customer who was involved in creating a service starts to consider it their own. I forgive minor bugs, because I know they get fixed. I do not leave silently, I write to you, because I believe I will be heard. And I defend your product in the comments more fiercely than your PR person, because I am defending something of mine. Calculate what one such brand advocate is worth. And a panel can hold ten of us.

When you cannot do without it: the AI makes decisions about me. Scoring, personalized pricing, approvals and rejections, claims handling. Every case where a mistake hits my wallet, not just your statistics. For such services, Chapter 3 has a special regime in store: a passport with a direct warning that a machine makes the decisions, and a receipt explaining why the decision came out the way it did.

How to choose your step

Ask yourself one question: what happens if the AI makes a mistake? If I shrug and rephrase, step one will do. If I get angry and leave, you need step two. If I seriously lose money, time or nerve, climb to step three and do not skimp.

And one small observation to finish. Companies that reach the third step almost never climb down. Not because they are stuck. Because they have tasted it: permanent access to a living customer turns out to be cheaper and more accurate than any analytics bought from outside. Trust a man who spent twenty years on both sides of these studies.

Which step are you on right now? If the honest answer is "none," do not worry. Read on. The next section is about how to ask correctly. Because asking is a skill too: ask the wrong way, and I will lie to you. Here is how to avoid that.

1.3. How to Ask So I Don't Lie

I must warn you about the main trap. If you ask me point-blank, "Will you like our new AI assistant?", I will answer, "Yeah, probably." And I will be lying.

Not out of malice. I feel awkward upsetting a person who is glowing while describing their project. Everyone does this. Even your mom will say the idea is wonderful. Especially your mom. There is a whole book about this called The Mom Test³, and its core idea is simple: people lie to you out of politeness whenever you ask in a way that makes lying easier than disappointing you.

Remember three rules for talking to me.

First: talk to me about my life, not about your idea.

Second: ask about specific things from my past, not my opinions about the future. Otherwise I will sincerely tell you I will use your bot every day, and then never open it once.

Third: talk less, listen more. The more you talk in the meeting, the less you learn about me.

Compare and feel the difference:

Bad: "Would you use a chatbot for booking?" Good: "Tell me how you booked with us last time. How many attempts? What annoyed you?"

Bad: "Which assistant features do you need?" Good: "What did you do when you couldn't reach us by phone? Walk me through it."

Bad: "How do you like our idea?" Good: "What did you overpay for, wait for, or drive somewhere for last month when you didn't want to?"

Bad: "How much would you pay for such a service?" Good: "How much did this problem cost you last time? In money, hours, nerves?"

Memories of my real suffering are worth more than fantasies about your bright future.

And there are three kinds of garbage you will hear from me that must be thrown away, not carried into a report.

Compliments: "great idea, well done." Pleasant, useless, forget it.

Fluff: "I would definitely," "I always," "I never." Anything that starts with "I would" is worth nothing.

Ideas: "you should also add voice input and a dark theme."

My ideas are not a development backlog. They are a signal to dig: "Why do you want voice input? Tell me about the last time you needed it." Sometimes a real pain hides behind an idea, but it is never where I am pointing.

Now the arsenal. Five tools, from cheap to expensive. Each has its strength, and none replaces the others.

The log. I already mentioned logs. Start here: it is free, and you already have it. My preferences are visible in my traces: at which step I abandon the dialog, which question I ask ten times, at what hour I storm into the chat, after which answer I leave and never return. Behavior does not lie. I may write in a survey that

everything is fine, but if eight out of ten people like me abandon checkout at step three, then step three is where your trouble lives.

The method's weakness: logs tell you where, but stay silent about why. For the why, move to the next tools.

My complaints and inquiries. Free gold sitting in your desk drawer. Export the last three months of customer messages: what people complain about, what they ask ten times, where conversations break off. The top 10 recurring questions are a ready-made brief for your AI, written by me personally, free of charge and without consultants. By the way, you can have the same AI analyze the export. Here is the prompt:

Here is an export of customer inquiries for the last __ months: [paste].

1. Group the inquiries by topic and count the frequency of each.
2. Write out the top 10 recurring questions verbatim, exactly as customers phrase them.
3. Flag where customers express irritation and what caused it.
4. Highlight questions that took more than a day to answer.
5. Compile a list of the customer's 5 main pains in their own words.

Interviews. Ten conversations of twenty minutes each. Expensive in time, but this is the only tool that answers the question "why." In a conversation I will tell you things you would never think to ask: that before calling you I call your competitor, that my daughter fills out the form on your site because I cannot read the small print, that the word "offer agreement" in your bot scares me. No survey will catch this, because a survey has no field for what you did not anticipate. The rules of the conversation are above: my life, my past, listen.

Surveys. A tool for validating with numbers what you already heard in interviews. First hypotheses, then the questionnaire, not the other way around. Keep it short: five to seven questions, more than that I will not endure. And I beg you, no "rate your satisfaction with our service level on a ten-point scale." Normal questions sound like this: how do you usually contact us, how many times did you have to repeat your question, what should we fix first. The strength: numbers you can lean on when deciding. The weakness: a survey answers only the questions you thought to ask.

Observation. Watching a live customer in their natural habitat is quite fascinating. This is the consumer test from step two. Sit me in front of your prototype and stay silent. Do not prompt me, do not make excuses, do not say "oh, that part isn't ready yet." Just note where I got stuck, what I muttered, and at what moment I reached for my phone to call the old-fashioned way. Painful? Very. More useful than any report? Absolutely. One hour of such observation saves a month of rework.

Where to start if money is short and time is shorter: logs and complaints today, they are free. Then ten interviews. That is enough to avoid missing badly with your first AI service. Surveys and tests can wait until there is something to validate and something to test.

Chapter 1 Checklist. Check Yourself

No.	Question	☐
1	We know the three groups of customer needs: functional, informational, emotional, and the nine dimensions within them	☐
2	We exported and analyzed complaints and inquiries from the last three months	☐
3	We have the top 10 customer questions, recorded in their own words	☐

4	We looked into the logs: where customers abandon the dialog and after what they leave silently	☐
5	We conducted at least ten interviews by the rules: about the customer's life, about the past, listening	☐
6	We separated facts from compliments, fluff and ideas	☐
7	We chose our step of collaboration: ask, let them try, involve	☐
8	A list of the customer's five main pains, in their own words, hangs in front of the team	☐

Eight items. Closed them all? Congratulations, you now know more about me than ninety percent of your competitors know about their customers. And now let's find out something unpleasant: whether you need AI at all. Don't rush to take offense. The next chapter may save you a great deal of money. We will be counting by my yardstick.

³ Rob Fitzpatrick, The Mom Test, 2013.

Chapter 2. Why Would Your Customer Need Your AI?

Allow me to ask a question no vendor will ever ask you. Everyone around is adopting AI, and so did you. Well done, congratulations. But what is in it for me?

Seriously. If you deployed AI for internal matters, to cut costs, replace staff, speed up paperwork, that is your business. Manage it as you wish. My opinion is not required.

But the moment you place AI between yourself and me, everything changes. Now, to get your service, I am obliged to talk to your bot. Not "allowed to," mind you, but obliged: you often leave me no other door. And since you made that decision for me, at least answer one question: will this make things better or worse for me? Not for you. For me.

The question is not rhetorical. It has a precise numerical answer, and in this chapter I will give it to you. Then I will show you what to check before any automation, so you do not end up accelerating nonsense. And at the end I will reveal a secret: why your dashboard shows three seconds while I have been fuming for two hours. We will be counting by my yardstick. Fair warning: my yardstick is harsh. But honest.

2.1. My Three-Number Test

Let me start from a distance. For many years the World Bank compared countries using the Doing Business methodology.⁴ The idea was brilliantly simple: look at the state through the eyes of an entrepreneur. Take a government service, say, opening a small firm, getting a construction permit or connecting electricity, and count three things.

- How many procedures you have to go through.
- How many days it takes.
- How much it costs.

In one country a new business opened in a couple of steps and a single day. In another it took a dozen offices and a month of waiting. Countries fought for years over their positions in that ranking. Notice what was measured: exactly what a person feels first-hand. Steps, time, money. Nothing else concerns them.

Well, I evaluate your company exactly the same way. Every day, free of charge and without any ranking. I have the same three indicators, and I want you to start counting them before you sign the AI contract with a vendor.

The first indicator: my steps. How many actions I must perform from the moment "I need this" to the moment "I got it." Everything counts: finding where to write to you. Registering. Confirming my email. Filling out the form. Attaching a document. Calling because the form did not work. Repeating my question to an agent because the bot passed nothing along. Every handoff, every "now go to another section" is a step. Steps are good at hiding: inside your chat there seem to be two, but along the entire path from my need to the result there are eleven.

The second indicator: my time. How long the whole story lasts. Not your internal "average handling time," but my time: with the searching, the waiting, the repeat visits and the "your call is very important to us." The service does not begin when I write to your bot. It begins when my need arises, and it ends when the problem is solved and I have forgotten about it. Everything in between goes on my meter.

The third indicator: my costs. Everything I spend beyond the price of the service itself. Fees and surcharges. A paid support line. Taking a day off to visit your office because "this matter is handled in

person only." Printing, signing, scanning. And let's be honest, my nerves belong here too: they do not show up in your books, but they show up beautifully in my reviews of you.

What to do with this? Take one of your services, the most popular one. Walk its path as a customer, from need to result, and write down three columns of numbers: steps, time, costs. That is your "before."

Then take the AI project and honestly forecast the same three numbers "after." If the AI is already running, it is even simpler: walk the path again and compare. You already know where to get the data from Chapter 1: logs, complaints, ten interviews, watching a live customer. And a glimpse ahead: once your AI starts issuing receipts, these three numbers will collect themselves. The receipt records the start and end times, the list of actions and the amounts. The three-number test turns from a one-off expedition with a notebook into a habit that works at every interaction.

From here, arithmetic decides everything. There are exactly three possible outcomes.

Outcome one. My numbers went down. Seven steps became five. Three hours became ten minutes. "Go to the office" became "solved it with the bot over coffee." Congratulations, this is the case that justifies adopting AI in the first place. Launch boldly. I will say thank you, come back and bring friends. Notice that I never even asked which model you use or how many parameters it has. I do not care. My numbers went down. The rest is detail.

Outcome two. My numbers stayed the same, but yours improved. You cut support costs, unloaded your people, and I felt no difference: I used to solve my problem in a day, and I still do. Fine, that is legitimate. Then one request: do not tell me in your advertising that it was all done "for my convenience." I can count, and that kind of falseness is felt instantly.

Outcome three. Your numbers went down and mine went up. The most common and the saddest outcome. I used to call, and a polite living person solved my problem in five minutes. Now I circle a bot for fifteen minutes while it fails to understand the question, then wait for an agent, then tell the whole story again because the bot passed nothing along. Your report shows savings. My meter shows that I am paying for your savings with my own time. That will not do. I do not need that kind of AI, and I will find a way to let you know: with a review and with my wallet.

Write this down.

AI adoption is justified for the customer when it reduces their steps, time and costs.

At least one of the three numbers, without worsening the others. Everything else is either your internal optimization or my degradation for your benefit. The first is acceptable. The second I will not forgive.

And do not panic if the test shows something unpleasant. It is far cheaper to see it on paper now than in the reviews later. The next section explains what to do with the steps the test caught. Spoiler: not all of them deserve acceleration. Some deserve a funeral.

2.2. The AI Guillotine: A Presumption of Corporate Inefficiency

Every company accumulates ballast over the years. Forms nobody reads. Approvals that approve nothing. Reports nobody needs. Steps that exist for a single reason: "that's how we've always done it." You no longer notice this ballast. You are used to it. I notice it every day, because I waded through it to reach your service.

And here is the thought this whole section was written for. AI adoption is the best excuse for a general cleanup your company will get in a decade. You were going to rebuild your processes anyway. So seize the

moment: do not drop a bot on top of what you have now. Upgrade the company. I call this maneuver the AI guillotine: AI adoption as an occasion to improve my service, which is the whole point.

The essence of the guillotine is a flipped lens. Normally a process lives until someone proves it is unnecessary. And nobody rushes to prove it: every process has an owner, the owner has a habit, and the habit has an iron argument, "we've always done it this way."

The AI guillotine flips the burden of proof: from now on, every process is presumed unnecessary until proven otherwise. A presumption of inefficiency for your operations. Every form, every approval, every step walks up to the whiteboard and explains why it exists. Failed to explain? Under the blade. No mercy and no farewell speeches.

Everything is cut in two passes, and the order matters. First remove what should not exist at all: duplicated and obsolete demands on customers. Then honestly divide what remains between humans and AI. In that exact order, because you cannot automate ballast. Or rather, you can, but then you are paying money to make nonsense happen faster. And removing it later will cost twice as much: by then it will have a budget, an implementation report and a defender with a slide deck. And the survivors get a prize: every service that lives through the guillotine will receive its own AI service passport in Chapter 3. Ballast does not get passports.

The moral: the AI guillotine is a conversation about the meaning of the work, not about technology, and it must be held with the employees who deal with me directly. And there is no better time than now, while the window of opportunity is open.

So you do not think this is pretty theory, one number from my own practice. I personally took part in such a "cleanup" for a city government. We analyzed more than a hundred services and cut their combined delivery time by more than five hundred days, and six services were eliminated entirely.

Five hundred days sounds modest until you remember that a service is not received by one person. Even if each of those services is used by just a thousand people a year, the city now waits one and a half thousand years less. Every year. Fifteen hundred years of human life that used to be spent waiting and have now been returned to people. Feel the scale.

Now, why do I, your customer, root so passionately for your AI guillotine? I have skin in this game. Three kinds of skin, actually.

The first interest is obvious: my savings. Most of the steps that will not survive the guillotine are my steps. My forms, my waiting, my "please send the document again." You cut the ballast, and my life gets lighter.

The second interest is pragmatic. The money for bad processes and for quality services comes out of the same pocket: mine. Every coin you spend feeding ballast is a coin you did not invest in me: in speed, in people, in the very AI that actually solves my problems. You feed the chaos less, I get better quality.

The third interest is human. The guillotine cuts activities, not people. Free your employees from routine, and they will finally have time for what you cannot delegate to a machine: complex cases, attention, a live conversation, empathy. With me, by the way.

And I know how to be grateful. I repay care in the same coin: I return, I buy, I bring friends. That is how this works. In the next section I will show you why, even after the cleanup, your reports may keep telling you fairy tales: time on your stopwatch and time on mine run at different speeds.

2.3. Time Is Relative: Your Instant and My Eternity

To finish the chapter, a magic trick you will encounter so often that it deserves its own section.

Picture a management meeting. Your team shows a dashboard: after deploying the AI assistant, the average response time is three seconds. Three seconds! Applause, bonuses all around. Now the same story through my eyes.

Say I want to change my order. First I spent 5 minutes hunting for where to even write about it: in my account, in the cart, in the contacts. Found it. Then I spent 10 minutes explaining to the bot what I wanted. The bot cheerfully replied in three seconds about my order, but not about the change. I rephrased. It again replied in three seconds, about my order. I typed "agent." It offered to rate the conversation. After 15 minutes of waiting a live person joined, and I spent another 5 minutes telling them everything from scratch. Total: 40 minutes of my life. On your dashboard, my nightmare looks like a series of brilliant three-second responses.

Neither of us is lying. We simply hold different stopwatches. Yours starts when I write to the bot and stops when the bot replies. Mine starts when my need arises and stops when the problem is solved and forgotten. Everything in between lands on my bill: the searching, the waiting, the repetitions, the transfers, the "tell me from the beginning."

This is not an AI disease. It is the old disease of all services. A car service shop boasts: "the part swap takes twenty minutes." And for the customer: book a week ahead, leave work early, drive over, wait in line, wait for the procedure, wait for the paperwork, drive back. Twenty minutes of work cost the customer half a day of life. Three seconds of your AI mean 40 minutes on my planet. Interstellar in action.

Why this is dangerous in practice. The metric you pay bonuses for starts to govern behavior. Reward response speed, and you will get fast responses. Fast does not mean useful: the bot will lightning-quick answer the wrong question, instantly close dialogs and swiftly ship me off to "the help section." Your stopwatch will set records. Mine will spin like a taxi meter. And then someone at a meeting will wonder why customers are leaving despite such wonderful metrics.

So write these down:

Measure the customer's entire journey: time from need to resolution.

The whole path, from my first action to the solved problem. Not average response speed, but how long the entire story took.

The share of issues resolved on the first attempt. Truly resolved: I got what I wanted and did not come back with the same question. A closed dialog and a solved problem are different things. We have covered this.

Repeat contacts. How many times I came back with the same problem. Every repeat means that last time you reported a resolution that did not happen.

The number of transfers and repetitions. How many times I was passed from hand to hand and how many times I told my story from scratch. Every "tell me from the beginning" subtracts from my love for you.

How do you see your real time? There is one way, and it is free: walk the path yourself, or ask someone from your household. Do not ask your team. They will show you the dashboard. Take your phone, imagine you need to reschedule a delivery, return an item or move an appointment, and walk the entire path as an ordinary customer. With the button hunt, the bot, the waiting. Time it from "I need this" to "done." It is

the cheapest audit of your life and, I suspect, the most sobering. Incidentally, in Chapter 5 I will teach you to make this check regular and systematic.

For now, let's wrap up the chapter. You have learned to count AI's benefits by my yardstick, held a general cleanup and now own two stopwatches instead of one. Check yourself.

Chapter 2 Checklist. Check Yourself

No.	Question	☐
1	We counted the customer's three numbers: their steps, their time, their costs. Before adoption and after	☐
2	We honestly identified our outcome: better for the customer, the same for the customer, worse for the customer	☐
3	We ran the guillotine: removed unnecessary steps and processes before automation, not after	☐
4	We divided what remained between humans and AI, after listening to the people who do the work	☐
5	We measure the customer's time from need to resolution, not the bot's response speed	☐
6	We track repeat contacts and "tell me from the beginning"	☐
7	Someone from management walked the customer's path personally, stopwatch in hand	☐

Seven items. If they are all closed, you hold something rare: an honest answer to the question "why does our customer need our AI." Most companies have a vendor's slide deck instead of an answer.

Now that we have established your AI is genuinely useful to me, it is time to agree on the rules of the game. What your bot promises me, what I am entitled to demand of it, and where its authority ends. That is Chapter 3, and it is about the most underrated word in business: the word "standard."

⁴ The World Bank, Doing Business reports (2003-2021).

Chapter 3. Agreeing on the Rules

So, we have confirmed that your AI genuinely helps me: it cuts my steps, time and costs, and the ballast has been thrown out. Now, before I start using it, let's agree on the rules of the game.

Because right now, hand on heart, I am playing against the casino. Only you know what your bot can do. You decide how long I wait for a reply. When to call a human is determined by an algorithm I have never seen. And when something goes wrong, you shrug: "technology, what can you do." No. Since you have placed a machine between us, put down in writing what it owes me.

3.1. AI Policy: How Things Are Done Here

The first document is the most important one in its consequences. Before I talk to your bot, I want to know what you have allowed your AI to do when dealing with me. Not within one specific service: in principle, across the whole company. That is what the company's AI policy is: an open document that hangs in a visible place on your website and is titled, say, "Our AI Policy."

Note: one clear document with references to the many legal ones. Not thirty pages of user agreement where a lawyer hid everything important between the commas. Not a privacy policy that no living person has ever read to the end. And not ten different pages scattered across the site. I will not hunt them down and glue them together. One document, in human language, two or three pages.

What do I want to find in it? Answers to ten simple questions.

- Who answers: which legal entity, and who inside the company owns the AI direction.
- What your AI does: the list of your AI services.
- How you tell me about it: when and in what form your AI introduces itself.
- How my data is handled: my consent, the collection, processing and storage of data, whether you train your models on my conversations, and whether I can delete my data.
- Whether your AI makes decisions about me: and if so, how I can get an explanation.
- How to call a human: the contacts, the timing, and the promise that a human will answer.
- What happens if your AI makes a mistake, what guarantees you give, and how you fix mistakes.
- How you announce changes: the document's version, its date, what changed.
- What traces remain: what gets recorded and what receipt I get to keep.
- And where to complain: the contact, the response time, how the complaint is handled, and refunds.

And let me point something out right away, so you do not mistake this for my whims. Parts of items three and four are already enshrined in law, though not identically for everyone. From August 2, 2026, Article 50 of the European AI Act requires the provider of a system that interacts directly with a person to ensure clear disclosure of the AI from the start of the first contact, unless it is obvious. For synthetic content covered by the law, the provider ensures machine-readable labeling. A company deploying the system separately warns about emotion recognition and biometric categorization, and explicitly discloses deepfakes and certain texts on matters of public interest if they are published without proper human oversight. These are different obligations for different roles and scenarios, not a single requirement to put a visible label on any text AI has touched.

Personal data handling is governed by the GDPR and related rules regardless of whether the company uses AI. The legal basis, the purpose, data minimization, retention periods and the person's rights must be determined separately for each service.

Violating the transparency requirements can cost up to fifteen million euros or three percent of turnover. Using AI to infer emotions in the workplace and in education is separately prohibited, except where justified by medical or safety reasons; for prohibited practices the fine can reach thirty-five million euros or seven percent of turnover. Note: the percentage is calculated on global turnover, not European revenue. For a company with fifty million in revenue, one such fine can devour the entire year's profit. The giants have the fat to survive it. A mid-sized business does not.

And do not exhale with the thought "this is about big tech." Put someone else's bot on your website? You may turn out to be a company deploying AI, a deployer, but that does not automatically mean all of a provider's obligations apply to you. You need to determine your role, the system's purpose, the nature of the interaction, the data and the risk level. The law distinguishes between those who developed a system or placed it on the market and those who use it under their own responsibility; their sets of obligations differ. But hiding behind the vendor contract will not work: the company must understand which part of transparency and customer protection remains on its shoulders.

All of this policy is written once, not from scratch for every bot. From there a construction kit takes over: each service passport simply pulls the sections it needs from the policy and refines them for itself.

And here, by the way, is the promised saving: when the support email or the name of the responsible person changes, you edit one place, not a hundred prompts. Engineers call this a single source of truth. I call it common sense.

The company's AI policy must have a version and a date. A policy without a date is a monument, and I need a living document. Keep the old versions too: a conversation that happened in March is judged by the March policy, not by the one you rewrote in April. And third, my favorite: a passport may refine the policy, but it may never quietly make it worse for me. Promised a human within five minutes in the policy? Then not one of your services has the right to quietly turn that into five hours.

What do you get in return, besides my trust? A prosaic benefit: some of this information is already required of you by law, though in different documents and only for the applicable scenarios. The rest you adopt as a voluntary consumer standard. By gathering it all into one readable document and marking what is law and what is your additional guarantee, you get not a formal brush-off but a showcase of maturity you can show me, the regulator and your partners.

You can add a register of trust credentials to the AI policy: current licenses and certificates, adopted standards, quality marks, results of independent audits and links to public test reports. It is not a fourth core document but a regularly updated annex. The policy shows the overall picture, and the passport of a specific service references only the credentials that apply to it.

Policy ready? Excellent. Now let's go down one floor, to a specific service. Meet the AI service passport.

3.2. The AI Service Passport: The Standard Your AI Works By

In Chapter 1 we mapped out my customer needs. Now a trick question: how do you make sure your AI actually meets them? Every day, in every dialog, including the one that happens at three in the morning on a Sunday?

The answer: give every AI service a passport. The AI policy we just wrote answers for the company as a whole, and it is deliberately general. The passport brings it down to the territory of one specific service. Remember the construction kit from the previous section? Here it is in action: the passport pulls the sections it needs from the policy (data, escalation, guarantees, complaints) and adds its own: what exactly this bot can do, what it promises and where its boundaries lie.

The passport is not technical documentation. Which model you use, which prompts, how the integrations work: you may write it down, but as a rule I am not interested. That is your kitchen. The passport describes not the bot's insides but the rules by which I can receive a proper service. That is: what I am entitled to, what I am not, and what I will sue you over.

Here is what I, your customer, want to see in the passport of every one of your AI services:

- What the service is, who provides it, how, and to whom.
- What I will receive and by when. The result of the service and the time you promise to deliver it.
- What the bot can and cannot do. The full list of actions and, more importantly, the boundaries: it processes returns but does not settle payment disputes. Honest boundaries save my time.
- What data the bot will ask of me. An exhaustive list. Anything beyond the list, the bot has no right to request.
- The cost of the service. If the service is paid, the price is in the passport, not a surprise at the end.
- On what grounds I can be refused. A closed list. "Denied" without a reason from the list is not a refusal.
- When the bot must call a human. Pre-defined cases: I asked directly; the bot failed to understand twice; the case is atypical; the conversation turned into a conflict. Pre-defined, not "at the algorithm's discretion." This is my emergency exit, and I want to see it on the map.
- What gets recorded. Every bot decision is logged: what it decided, by what criterion, when. A black box, like on an airplane. If we argue, we pull the recording.
- A check before delivery. Before handing me the result, the bot verifies it against the passport. Not sure? It does not deliver. It calls a human. Better to wait a minute than to receive confident nonsense.
- Where to complain and how fast you will answer. The channel, the deadline, the form of the response.

Look closely at this list: it is my needs from Chapter 1, translated into rules. The functional ones became deadlines and boundaries, the informational ones became lists and grounds, the emotional ones became the emergency exit to a human. The circle has closed: you asked what I need and wrote it into the document the machine runs on.

Now the rules that make a passport real rather than decorative.

Rule one: the passport is public. It hangs where I will find it: next to the chat, on the service page. Written in human language, briefly. Not "your request is being processed," but "we are looking into it and will reply within a day." Every piece of officialese in the text reads to me as "we can't be bothered to talk to you like humans." A passport buried in a forty-page user agreement exists legally, but not humanly.

Rule two: word equals deed. A gap between the passport and the bot's actual behavior is not a "documentation inaccuracy." It is misleading the customer, with everything that follows. Wrote "a human joins within two minutes"? Kindly make it two.

Rule three, my favorite: you answer for the AI. The phrase "it was the algorithm's error, we are not to blame" does not work. The bot quoted a price? That is your price. The bot promised a deadline? That is your promise. A company may not hide behind its AI, just as it may not hide behind an intern: you hired it, you answer for it.

And a fourth detail, small in appearance and important in a dispute: the passport lives in time. It has a version, a date and a status: in testing, in force, suspended. Updated the model, changed the price, taught the bot something new? That is a new version of the passport, and the old one stays in the archive. Why

such scrupulousness, you will understand a few pages from now, when we get to the receipt: a conversation that happened in March will be judged by the March passport, and that protects us both.

A small glimpse of the future, so you understand where all this is heading. In a year or two it will not be me coming to you for a service, but my personal AI assistant: to compare prices, place an order, ask questions. Your beautiful website means nothing to it. It needs the AI service passport in machine-readable form: what you do, what you require, what it costs, where the machine-readable service cards are, whether you issue a receipt, and where my agent should send its requests. Whoever has a passport, my agent will strike a deal with in seconds. Whoever does not, it simply will not understand, and it will move on to your competitor. So the passport is not bureaucracy. It is your future language for talking to an entire new class of customers in the agent economy.

And for you personally the passport is convenient for one more reason: it is a source of public requirements and acceptance criteria. "Build us a good bot" is not a brief. An AI service passport can be attached to a vendor brief, and the promised result can be checked against it. But a full technical specification is broader: it covers architecture, models and tools, access rights, integrations, data formats, security, test scenarios and fallback modes. The passport answers "what is promised to the customer"; the technical specification answers "how to implement it reliably." As a bonus, disputes with me turn into arithmetic: promised in the passport, delivered, case closed.

So: the policy said how things are done at your company. The passport promised how one specific service will work. One question remains, and it is the most uncomfortable one: who will prove that in my particular conversation everything went exactly as promised? A promise without proof is advertising. So allow me to introduce the third document, my favorite. The AI receipt.

3.3. The AI Receipt: What Actually Happened

Allow me a short history lesson; it is worth it. When the consumer movement was just being born, it had no laws, no associations, no hotlines. It had one rule: always take the receipt.

That rule was taught from posters and screens, repeated like an incantation. Why such honor for a simple slip of paper? Because the receipt is the only door to justice. The entire majestic temple of consumer protection stands on that scrap of thermal paper.

Now look at what is happening in the world of AI. Your bot quoted me a price, promised a refund, calculated my insurance, booked my appointment. The conversation ended, the window closed. What do I have left in my hands? Nothing. At best a chat history that lives somewhere on your side, gets edited by you and disappears when it suits you. A hundred years of consumer progress, and I am standing at the counter empty-handed again. Only the counter is virtual now.

So my request is simple: have your AI issue me a receipt. Not a forty-screen transcript, but a document generated at the end of the interaction that stays with me. What goes in it?

- Who provided the service and to whom: the company, the AI assistant's name and version, the date, the start and end time, and the service recipient's details.
- What service was provided, with the full conversation in Markdown.
- What consents and customer data were received.
- The agent's actions: amounts, deadlines, next steps. Item by item, like lines on a cash register receipt.
- What the result of the service was.
- The cost of providing the service.
- Errors, downtime, rework, employee intervention.

- Under what policy all of this happened: the version of your passport and policy at the moment of the conversation.
- Where a claim can be filed.
- And a button in plain sight: "contact the company."

The AI receipt carries not a link to your policy but a snapshot of it: the full text of the exact edition in force at the moment of the conversation. The difference is fundamental. A link leads to a website that can be rewritten tomorrow. The receipt is a photograph of our agreements, and it cannot be retouched after the fact. Remember me droning on about versions and dates in the policy and the passport? This is where it fires.

The AI receipt must exist in two identical copies. One with me, one with you, both generated automatically from the same authoritative package of events. But two files are not enough if either one can be quietly edited. So every receipt gets a unique number, a server timestamp, the versions of the policy and the passport, and a link where its authenticity can be verified. For significant interactions, add a hash or an electronic signature; a correction is issued only as a new version, and the previous one stays in the history. Now watch the beauty of it: the dispute over "what did your bot actually say" dies before it starts. No need to trust my screenshots or your word about the logs: we compare the verifiable copies. The receipt protects me from you and you from me at the same time, and I honestly do not know who needs it more.

When should an AI receipt be issued? Common sense, no fanaticism. I asked the bot for your opening hours? No receipt needed, don't be ridiculous. But the moment the AI does something substantial: places an order, runs a calculation, charges money, accepts my documents, gives a recommendation my wallet depends on, the receipt is generated automatically. Not "on customer request, if they think to ask," but by itself, always, automatically. If I want it, I open it; if not, let it sit on your side.

And if you think the AI receipt is my fantasy, look at where legislation is heading. In Quebec, Law 25 requires telling a person when a decision is based solely on automated processing, and grants the right to request information about the data used and the main factors of the decision, to submit observations and to have the decision reviewed by a human. That supports the logic of the receipt, though the law itself does not call such a package a receipt. In Utah, permanent clear disclosure of generative AI's nature can provide a limited safe harbor from a specific enforcement action for violating the disclosure requirements; it is not a universal release from liability. In the EU, logging obligations depend on the role and mostly concern high-risk systems, and the GDPR grants the right to a copy of personal data, not of all internal technical records. So the AI receipt is not some supposedly ready-made universal legal obligation, but a voluntary consumer standard that ties several legal trends together and immediately gives both sides proof of the interaction.⁵

I hear the objection: paperwork, where would a small business find the resources. I answer with pleasure. Who will write this receipt? Nobody, by hand. The information system generates it automatically from authoritative records: the event number and time, the dialog, the actions, the consents, the result, and the versions of the policy and passport in force. Generative AI can prepare a human-friendly summary, but it does not invent the canonical facts or reconstruct them from memory. They are pulled from the system like lines on a cash register receipt. So the document appears in a second, and the verifiable data in it matches what the company has stored on its side. Handing the customer such a copy is a cheap gesture. And it looks like a revolution.

One final touch. The receipt turns your support service from a propaganda bureau into a service. My complaint no longer starts with a half-hour retelling of "so yesterday your bot told me." It starts with a

receipt number. Everything is already inside: who, when, what was promised, what went wrong. You solve the problem in minutes, I write a glowing review, everyone hugs. Curtain.

Policy, passport, receipt. How things are done, what was promised, what happened. The great trinity, assembled.

3.4. Let Me Show You: Policy, Passport, Receipt

Theory without a sample is dead, so let me show the whole trio live, on one example: a language school, with a chat bot on the website answering questions and helping put together a trial lesson request. First the school's policy, then its bot's passport, then the receipt of one real conversation. Notice how the documents nest inside one another: the passport references the policy, and the receipt carries a snapshot of both.

OUR AI POLICY. Language school [name]. Version 1.2 of 01.07.2026

Who answers: [legal entity], owner of the AI direction: [name], write to: [email].

Where AI works here: a chat bot on the website and in messengers (consultations and putting together an enrollment request). The bot's passport: [link].

How the AI introduces itself: in the first message, always. Our bot does not pretend to be human.

Your data: name, contact and topic of the inquiry; kept for up to 12 months to handle the inquiry and a possible claim; we do not train our models on the dialogs; we pass data only to the processors listed at [link] that are needed to run the chat, the CRM and communications; we do not sell it or allow it to be used for their own purposes. On your request we will delete your data earlier.

Decisions about you: the bot makes none. Our prices and discounts are the same for everyone, and enrollment is confirmed by a staff member.

How to call a human: the button in the chat, the word "operator" or a call to [phone]. During working hours a human replies within 15 minutes.

If the AI made a mistake: the bot's mistake is our mistake. We fix it free of charge and honor what the bot promised.

About changes: a new version of this document is published right here, with a date; we keep the old versions.

What traces remain: the dialog is recorded; after a trial lesson request is submitted, the system automatically generates an interaction receipt.

Where to complain: [email], reply within 2 business days. No reply? [external ADR channel].

AI SERVICE PASSPORT: course information and trial lesson requests

1. General provisions. Service: information about courses, schedules and prices, and putting together a trial lesson request. AI agent autonomy level: information and creating the request; enrollment is confirmed by a staff member, payment is handled by a staff member. Who may apply: any visitor to the site or messengers. Result: an accurate answer, a request created where needed, and a clear next step. Introduction: the assistant opens with "I am the school's AI assistant." Feedback: a "Complain" button in the chat, a human reply within one business day.

2. The standard: what we promise. Covers: the list of courses and formats; schedules; prices and discounts; how to enroll and what happens next. Channels: website chat, messengers, email. Timing: instant replies in chat, waiting no longer than a minute; email within one business day. What it does not do: does not judge your language level "by eye," does not give pedagogical advice, does not discuss other clients. Grounds for refusal: the question is not about the school (we will point you in the right direction); requests for other people's data. Price: free. Quality: answers strictly from the approved standard, nothing "improvised"; the question is closed in full; handover to a human preserves the context. Data: only your name, contact and language level, nothing beyond that is requested.

3. The AI agent's procedure. Steps: introduce itself as AI, clarify the question, answer strictly by the standard, offer a next step and, if the client wishes, put together a trial lesson request. Must call a staff member if: the client asked for a human; it failed to understand twice; enrollment needs confirming, payment needs taking or a refund needs processing; the conversation turned into a conflict. May never: invent things absent from the standard; quote prices or dates not in the standard; confirm enrollment or payment on a staff member's behalf; give legal or medical opinions; request data beyond the necessary. Logging: every dialog is recorded: question, answer, whether a request was created, whether it was handed to a human, the client's rating. Check before delivery: the answer is verified against the standard; when in doubt, the bot does not answer on its own but calls a staff member.

Machine-readable card (for customers' AI agents):

```
service: course information
provider: language school [name]
description: courses, schedule, prices, trial lesson requests
channels: chat, messenger, email
price: free
response_time: instant; email <= 1 business day
languages: en, [your languages]
disclosure: service is provided by an AI assistant
escalation: enrollment confirmation, payment, conflict, on client request
contact: [link and phone]
issues_receipt: yes
updated: [date], version 3.1
```

INTERACTION RECEIPT No. 2026-0718-0142. Language school [name]

Date: 18.07.2026, 21:42-21:51. Channel: website chat.

Assistant: the school's AI assistant, version 3.1.

Introduced itself as AI: yes, in the first message.

In force: policy v1.2 of 01.07.2026, bot passport v2.0 of 15.06.2026 (full texts attached to the receipt).

You asked: an evening English course for level B1, starting in September.

The bot understood: find a B1 group, weekdays after 18:00.

The bot did and promised:

- offered the Tue/Thu 19:00 group, teacher [name], starting 02.09;
- quoted the price: 96 euros per month, no discounts for this group;
- put together a request for a free trial lesson on 26.07 at 19:00;
- confirmation of the lesson and enrollment will be done by a staff member, who will contact you by 21.07.

Human involvement: required to confirm the trial lesson and enrollment; the request was handed to staff member [name].

Your data in this conversation: name, phone, language level.

Result: the trial lesson request was created and handed to a staff member.

Cost of the service: free.

Errors and failures: none recorded.

Full conversation: attached to the receipt (Markdown).

Consent to be contacted: given at 21:49.

Disagree with something? [Contact the company] [Dispute] Claim by this receipt number: [email].

Verify this receipt's authenticity: [link].

Service owner: [name, position].

And that is all. Three documents, three pages of text, but look how many of my requirements they close at once. I know how your company treats AI and my data. I know it is a machine talking to me, what it can do and what it cannot. I know the quoted price is real, because improvisation is forbidden to the bot in

writing. I know my emergency exit: two misunderstandings and I am handed to a human instead of being run in circles. And a receipt stays in my hands, so a dispute, if one arises, will take a minute, not a month.

Notice four things.

First: the procedure section does not need to be pasted into your agent verbatim as its system prompt. The passport's requirements are translated into a controlled operating policy: system instructions, machine rules, tool permissions, checks and tests. The passport remains a public promise to the customer, and the operating policy makes sure it is kept. They come from the same approved source and change in sync, but they are not the same document. Word equals deed not because the text was copied into a prompt, but because the promise has been turned into verifiable constraints on the system.

Second: the documents nest inside one another like matryoshka dolls. The passport references the policy and may not quietly worsen it. The receipt carries a snapshot of both the policy and the passport in the version that was in force at the moment of the conversation. That is why the trio cannot be torn apart: pull on the receipt, and you pull out the entire chain of agreements.

Third: the machine-readable card. Remember I said my personal AI assistant will be visiting you soon? The card is a shop window for it, and the line "issues a receipt: yes" will be the same mark of trust for my agent as a cash register on the counter is for me. The fuller the card, the more willingly other people's agents will recommend you to their owners. Free advertising in a channel your competitors have not even heard of yet.

Fourth: the passport has an owner with a first and last name. Not "the team," not "the department," but a person who answers. To me, your demanding customer, this matters more than everything else.

Where to start? The order is simple, and it is already built into the trio itself. The policy is written once for the whole company: that is one evening of work with this brochure in hand. Then take one service, the simplest one, like this school example, and assemble its passport from the ready-made policy. Switch on the receipt. Launch, watch the logs, adjust. Then take on the next service: the policy is already done, only the passport remains. In a couple of months you will have something almost nobody on the market has: an AI that works by the rules, answers for its words, and leaves a trace both sides trust.

Chapter 3 Checklist. Check Yourself

No.	Question	☐
1	The company has public AI rules: one open document answering the ten questions, with a version and a date	☐
2	Every customer-facing AI service has a passport (general provisions, standard, agent procedure) assembled from the rule modules	☐
3	The passport is published where the customer will find it, written in human language, and fixes: what the bot can and cannot do, deadlines, data, price, a closed list of grounds for refusal	☐
4	The cases when the bot must call a human are spelled out: customer's request, two misunderstandings, money, conflict	☐
5	The bot introduces itself as AI at the start of the conversation; generated content and deepfakes are labeled (Article 50, deadline August 2, 2026)	☐
6	Employee emotion recognition is switched off; customers are warned about tone analysis	☐
7	Every dialog is logged, and for a significant service the customer automatically receives an interaction receipt with a snapshot of the rules and passport in force	☐

8	The rules and every passport have an owner with a first and last name, and a gap between the document and the bot's behavior is treated as an emergency, not a "nuance"	☐
9	The machine-readable service card is published for customers' AI agents, and it states: a receipt is issued	☐

⁵ Quebec (Canada): Law 25, art. 12.1, notification of an exclusively automated decision, information about the data used and the main factors, the opportunity to submit observations and obtain a human review. Utah (USA): Utah Code Title 13, Chapter 77, disclosure requirements and a limited safe harbor from an enforcement action under the relevant section. EU: AI Act art. 12, 19 and 26(6), role-based logging obligations for high-risk systems; GDPR art. 15(3), the right to a copy of the personal data processed, not of the entire technical log.

Chapter 4. Talk to Me Like a Human Being

You have built the AI agent, created the policy, the passport, the receipts. Now let's talk about how you present all of this. To me, your customer. This stage is critical: present it wrong, and I walk. And by the way, have you forgotten your employees? Don't be surprised: your team also needs the AI news delivered properly, to defuse their anxiety and potential sabotage, so that I, once again, don't end up the casualty.

Let's start with the foundation. There is no point informing anyone if you do it in a language that needs translation. Let me explain about clarity. It is a sore spot.

4.1. Speak Plainly

Open your bot's chat history with customers and look for phrases like "your request is being processed," "the provided data is undergoing verification," "under the current terms of service." Found them? Your AI speaks officialese, a dead language no living person uses. I write to the bot "where is my order," and it answers me like a bureaucrat replying to a complaint. It feels as if I hadn't tried to buy something but had violated something. This is genuinely infuriating.

Officialese used to have an excuse: lawyers wrote the texts, to be safe, and rewriting them was expensive. Your bot has no such excuse. A language model can speak simply, warmly and clearly. It is literally its native talent. If your AI talks like a tax inspector, someone set it up that way or was too lazy to set it up otherwise. It was a choice, not a limitation of the technology. And when I read "your request is being processed," I feel that choice on my skin: you couldn't be bothered.

The rules of human language are simple. Here are the six main ones:

Short sentences. One thought per sentence. If a phrase has three commas and two "whiches," chop it up.

Human words. Not "effectuate payment" but "pay." Not "in the event that questions arise" but "if anything is unclear." Not "functionality" but "what it can do."

Active voice. Not "a refund will be issued" but "we will return your money." In the first version the money returns itself, mysteriously, and there is no one to ask. In the second there is a "we," and that "we" answers.

Lead with the point. The answer to my question in the first line, details below. Do not make me read three paragraphs of preamble to learn "yes" or "no."

Concrete numbers. Not "shortly" but "by the end of the day." Not "a fee may apply" but "a 2 percent fee." Fog, I remind you, I do not forgive.

No internal jargon. Your "tickets," "escalations" and departmental abbreviations stay in the department. I have an order, a question and a desire to get an answer.

For clarity, before and after the translation into plain language:

Before	After
"Your request has been registered and will be reviewed within the established timeframe"	"We got your question, we will reply by 6 pm"
"To effectuate the refund of funds it is necessary to provide payment details"	"To return your money, send us your card number"
"At the present moment the given item is absent from"	"This item is out of stock, back on Friday. Want me to"

the warehouse"	hold one?"
----------------	------------

Feel it? In the second version I am talking to someone who cares, even if it is a machine.

The test is simple: read the bot's reply out loud, as if talking to a neighbor. If it cannot be said out loud, rewrite it. And to avoid rewriting by hand, give the job to the bot itself. Here is the prompt:

Here are our AI assistant's standard replies: [paste].
 Rewrite each one by the rules of human language:
 short sentences, active voice ("we will return", not "will be returned"), the point in the first line, concrete dates and numbers,
 no officialese or internal jargon. Tone: a friendly adult,
 no baby talk. Output a table: before / after.

Half an hour of work, and I will notice the effect from the first message. And here is the important part: plain language is not "dumbing it down," as some experts think. It is respect in its most practical form: you spend your effort to save mine. And it is in this very language that you will need to write everything covered in the next two sections: the story for me and the story for your team.

4.2. AI as an Occasion to Care About Me

Let me share a secret that will save you marketing budget. I, your customer, am deeply indifferent to your technology. Neural networks, models, agents: those are your toys. But I care to the point of trembling about how you treat me. Which is why one and the same event, "we have adopted AI," can be presented in two ways. The first: quietly install the bot and wait until I stumble upon it and draw my own conclusions. The second: turn the launch into a vivid gesture of care for me. Guess which way most companies choose. And guess which one pays off.

Here is what care sounds like. Not "we have implemented an innovative AI solution to optimize customer service." That is about you and your pride. But like this: "Dear customers, you are the most important people to us, so we built an assistant that answers you within a minute at any time of day. The live team hasn't gone anywhere: one click, and you are talking to a person. Our staff member [name] is responsible for the assistant. Write to them if anything is off. And send us ideas for improvement: we really do read them." The same event. The difference is who stands at the center: you or me. When it is me, I feel it, and I remember.

There are more channels for this story than you might think. Pick by your audience:

A letter to customers: "we adopted AI with care for you," two paragraphs, warm and without pathos.

A social media post in the same tone.

A short announcement page on the website.

A notice right in the chat: "An AI assistant is replying. Need a human? Tap here." This one, by the way, is also a legal requirement. Remember Chapter 3.

A client cheat sheet: how to use the assistant, three steps.

An FAQ with honest answers: is it really AI? is my data safe? how do I reach a human?

And if you have a physical location, hang a sign at the entrance: "AI works here: it helps faster, a human is always nearby."

Visitors smile, take photos and, most importantly, are not startled when a machine answers them. Each of these formats takes an evening of work. Together they turn a technical event into a story about how you thought of me.

By the way, a gift, so you cannot say I only make demands. Here are eight ready-made "AI works here" stickers, drawn for my client communication kit. Take them: put one on your website, in your chat widget, on the office door, in your newsletter footer. No licenses, no small print. The full editable kit with all the templates lives at aibusiness.vc.



Now further. I received your letter and came to your website. What should I find there?

You already know the answer: the AI policy from Chapter 3. That same single document about all your AI, in human language, with a version and a date. Here it plays a new role: not a legal backbone but a showcase of care. The letter invited me over; the policy greets me at the doorstep.

And let me remind you of the engineering benefit your own developer will thank you for: everything changeable (names, contacts, prices, hours) lives in the policy in a single copy, the bots pull it from there, and the company's word, outside and inside, matches by construction. Change it in one place, and it changes everywhere.

The same story awaits your bots. Today you have one bot, and you have calmly sewn into its prompt the name of the responsible person, the support email and the working hours. In two years you will have twenty bots, maybe a hundred. And then the employee quits, the email changes, the office moves. What, rewrite a hundred prompts? Why. All the information that can change: names, contacts, hours, links, prices, lives in one place, in that same public AI document. The bots do not store it. They pull it from there. Change it in one place, and it changes everywhere. Your document for customers and your reference for bots turn out to be the same file, and the company's word, outside and inside, matches by construction.

And one more reason to build a single source of truth, the most far-reaching one. Soon it will not only be me visiting your site but my personal AI agent: I will task it with finding the best deal, asking questions and preparing the purchase. Your beautiful design means nothing to it. It does not look at pictures. The first thing it will look for is who it can talk to here in its own synthetic tongue. It will want to know which chat bots you have, where the machine-readable service cards are, whether you issue a receipt, and where its requests should go. For it, the AI policy is not an "about us" page but the entry point for negotiations. Whoever has such a document, my agent will close a deal with in seconds. Whoever does not, it will honestly skip and move on to your competitor. So you are writing this document not only for today's me, but for tomorrow's stream of customers brought by their agents.

A beautiful structure has emerged. Appreciate it: the passport describes one specific service, the single source of truth describes the company. I learn about your AI before meeting it, communicate with it by clear rules, and always know the information is accurate and current. One audience remains, the one almost everyone forgets during AI adoption, and wrongly so: the people who work for you, your employees. It is your business, of course, but if something goes wrong with the AI, it is them I will be switched to, and them I will yell at. So they too deserve a warning. And herein lies the main value of this guide.

4.3. People, Not "Personnel Costs"

Your people serve me, and I feel their fear through the counter as clearly as I feel officialese in the chat.

Here is a scene playing out in thousands of companies right now. A bot was deployed. Nobody really explained anything to the staff, but everyone has read the news about "AI will replace millions of jobs." So I call, and a live manager confides: "Don't bother with that bot, it's useless. Let me sort you out like a human being." Recognize it? That is not a bad employee. That is a frightened employee defending their job the only way they can: by quietly sinking your bot. And I am the one standing between two fires, and it is not cozy here.

Now let's talk honestly about the root of that fear, which is to say, about you, dear executive. I know what crossed your mind when you heard about AI. Payroll is the fattest line of expenses, and the thought of

"replacing everyone with synthetic bots" is tempting to the point of trembling. All the more so because the report looks superb to investors: costs cut, well done, bonus for the board.

So allow your demanding customer to say something rude about this strategy. Rude, but arithmetically precise. You can only cut costs down to zero. You can grow endlessly. The agenda of "fire everyone to save money" is a false agenda, and here is the question that kills it: what will you do with the money you saved? Put it on deposit? Saving is not a strategy. It is a one-time trick. Growth is a strategy.

Look at the real trap of thinking here. Most executives see AI as a way to digitize their current level: do the same things, but cheaper. Few ask the next question: what level could we move up to thanks to AI? What new services could we offer customers that used to be beyond our reach: not enough hands, time, money?

Remember the AI guillotine from Chapter 2: we cut the fat not for the sake of saving as such. We were clearing space and freeing resources for growth. AI gives you processes that used to be fundamentally inaccessible. It is a unique chance to start doing what you have never done. Scary? Yes. But trading that chance for "we laid off three operators" is like buying an airplane to haul potatoes to the market next door.

And now about your people. Listen carefully, because I am speaking as the customer. Your employees know me better than any bot. They have worked with me for years. They know my moods, my emotions, my quirks that are written down nowhere. They remember what I look for, what I complain about, what my business is, when my season comes and what I can be offered. The bot has none of this knowledge. It is not in the CRM. It is in heads and fingertips. In the atmosphere that forms at the moment of a deal.

Fire the people, and you throw away the most valuable knowledge in the company: the knowledge of me, your most important person. The right scheme is exactly the reverse: routine to the AI, the customer to the human. Let the machine fill out forms and answer the routine, while a living person guides me, watches over me, talks to me warmly and offers what a bot would never think of.

And one more thing that changed quietly and radically: every one of your employees can now be a builder. Assembling an AI agent today is easier than mastering a spreadsheet with macros. Someone on your customer team has surely been sitting on an idea for five years: "we really should also do this, but there's no one to do it." Well, now there is someone: her, with AI in her hands. Ask your team not "who do we cut" but "what new things can we now offer customers." The difference in the question is the difference in the company's fate.

How to tell all this to the team, short and to the point:

First: formulate the principle and repeat it. We are cutting activities, not employees. Routine goes to the machine, people move to more interesting work closer to the customer. Saying it once is not enough. People believe deeds: the first employee freed from routine who receives a new role, not a termination letter, will convince better than a hundred town halls.

Second: train them. What the bot may do, what it may not, which data must never be uploaded, how to check AI answers, how to hand a customer from the bot to a human so the customer does not retell everything. An hour at the start, a short refresher every quarter, onboarding for newcomers. And by the way, this is no longer just my advice: since February 2025, European law directly obliges companies to ensure adequate AI literacy⁶ of staff working with AI. While everyone was discussing the 2026 deadline, this obligation had already arrived.

Third: open an idea channel. Let anyone propose: "let's have the AI do this." Half the ideas will turn out to be gold, because your people see what neither you nor consultants see: the customer's real life.

Fourth: appoint an owner for the AI household. The company needs at least one person for whom this is written into the job description, not into their enthusiasm: collecting employees' ideas and turning the best into working bots, writing briefs for vendors and accepting their work, maintaining the service passports and that single source of truth, training the team, following the news of models and laws, which, as you have noticed, update more often than price lists. In a small company it is one employee with dedicated hours, in a large one a whole department. Call it the AI competence center. And I, the customer, benefit directly: when something goes wrong, I know the company has a specific person who will learn about it and fix it, not "the thing sort of runs itself."

Chapter 4 Checklist. Check Yourself

No.	Question	☐
1	The bot's replies pass the "read it aloud to a neighbor" test: no officialese, the point in the first line, concrete deadlines	☐
2	Customers were told about the AI adoption as an act of care: what it gives them, a human nearby, a named person responsible	☐
3	Clear channels were used: a letter, a post, a notice in the chat, a cheat sheet, an FAQ	☐
4	The website has one public document about all the company's AI: services, boundaries, data, responsible person, where to complain, last updated	☐
5	Changeable information (contacts, names, prices) lives in a single source of truth, and the bots pull it from there instead of storing it in prompts	☐
6	The team has been told the principle "we cut activities, not employees," and it is backed by deeds	☐
7	Employees are trained: the rules, forbidden data, checking answers, handing a customer to a human (a legal requirement since February 2025)	☐
8	An idea channel is running: what else can we offer customers with AI	☐
9	A head of the AI competence center is appointed: an employee responsible for ideas, bots, passports, training and news	☐

In sum: it takes plain language, an honest story and respect for two groups of people: those who buy from you and those who work for you. And now the final chord of this guide: let's check how all this beauty actually works. Through my eyes. Chapter five: complaints, mystery shoppers and rules that live.

⁶ EU AI Act, Article 4 (AI literacy), in force since February 2, 2025.

Chapter 5. How to Evaluate AI

So, you did everything right: asked me, counted the benefits, wrote the AI service passport, presented it to everyone like a human being. The bot works, the dashboard is green, the team is happy. One last question remains, the trickiest one: how do you know it works?

Just do not show me the dashboard again. We established back in Chapter 2 that your stopwatch and mine run differently, and that "dialog closed" does not mean "problem solved." Internal metrics tell you how the machine labors. The truth about service quality lives on my side of the counter. This chapter is about building an evaluation that looks through my eyes: without a budget for sociologists, without reports for the drawer, and without self-deception.

5.1. Build an Evaluation System, Not a Collection of Likes

Let me start with the main rule, which will save you a great deal of time and money. Before collecting any number, any review, any rating, ask one question: what decision will we make based on this data? If there is no answer, do not collect it. Data without a decision is not analytics. It is stamp collecting. Ask yourself: what happens if we do not have this information? If the honest answer is "nothing," then there is no reason to gather it.

Hence the second rule, twin brother of the first: it only makes sense to evaluate what you control. If a parameter cannot be changed, measuring it is entertainment. You control the bot's response speed? Measure it. My mood on a Monday? Leave that to the astrologers.

Now, what exactly to evaluate. Everything is already in place. We built it in the previous chapters. It just needs assembling. The evaluation of an AI service stands on two pillars.

The first pillar: are the promises kept? You have a service passport, and a passport is a ruler. Promised a reply within a minute: what share of dialogs made it? Promised a human in one click: how many clicks did it actually take? Promised not to ask what you already know: does it re-ask? Every line of the passport becomes a verifiable indicator. That, by the way, is why I insisted so hard on measurable promises: you cannot put a ruler against the slogan "we strive for the best."

And with the arrival of the receipt, the ruler has gained a partner: the fact. The passport says how things should be. The receipt says how they were. Evaluating AI turns into a boring and beautiful procedure: take a month's stack of receipts, lay them against the passport, count the discrepancies. Every discrepancy is either a defect in the bot or an outdated passport, and both are curable.

The second pillar: did things get better for me? Remember my test from Chapter 2: steps, time, costs? It works after launch too. Plus my feeling: is the problem solved, will I come back, will I recommend you. The first pillar without the second is dangerous: you can flawlessly keep promises I do not need.

How to collect data without going broke. A small business does not need sociology with hundred-person focus groups. There is an approach called lean data⁷, built for verifying real impact without big budgets.

Four principles: ask the customers themselves, not experts; collect exactly as much as the decision requires; use simple free tools; ask briefly and often, not lengthily and once a year. A five-question survey sent on time beats an annual study lying in a drawer.

And what to ask me, so as not to annoy me or waste my time. Here is a five-question backbone as a gift. Adjust it to your needs:

Did you manage to solve your problem with the assistant? If not, what got in the way?

How much time did the whole thing take, from "I wanted" to "I got"?

Compare it with how things were before: simpler, the same, harder?

Would you recommend us to friends? Why?

What should we fix first?

Pay attention to the third question. It is my favorite. Evaluation always needs a comparison: with how things were, with a competitor, with expectations. "Did you like it?" is a question about nothing. People answer "yes" out of politeness. We covered that in Chapter 1. "Did it get simpler than before?" is a question with an honest answer.

What remains is to assemble a system out of this, meaning add a calendar and an owner.

Continuously, by itself: logs (where dialogs get abandoned, where people call for a human, where the bot re-asks) and a one-click micro-rating after each dialog.

Quarterly: a short survey of the base using those same five questions, plus a check of every promise against the passport.

Twice a year: the exercise described in the next section. Patience.

Every number, I repeat, must end in a decision: you saw that a third of dialogs break at step three, you fix step three and look again a month later. You saw it and sighed: consider that you never saw it.

That, essentially, is the whole infrastructure: the passport as a ruler, five questions as a thermometer, a calendar as discipline, and one rule on top: no decision, no measurement. And now the tastiest part. There is an evaluation method that I, your demanding customer, love most tenderly of all. So much that I am ready to take part in it personally. Next section: the mystery shopper.

5.2. The Mystery Shopper: Hire Me

A confession of love. Of all the evaluation methods I have seen in twenty years, this is my favorite. It is cheap, honest and hits dead center. It is called the test purchase, or mystery shopping, and if you take only one method from this whole chapter, take this one.

The essence is indecently simple: walk through your own service as an ordinary customer. Not as the executive who knows where every button is and what was meant. As a person off the street: anonymously, with a real need, from "I need this" to "I got it," writing down everything that happens along the way. The view not from behind the counter, but from my side of it. The difference, as you will discover, is enormous.

Why this method seems made for AI services. First, testing a bot requires neither agents nor a budget: no actors to hire, no branches to drive to. A laptop and half an hour will do. Second, you can repeat it every week if you like. The bot will not recognize the inspector and start trying harder. Third, and most important, a live walkthrough catches what logs will never show: dead ends with no exit, a tone that sets your teeth on edge, and that feeling of "I am nobody to them," which no metric records but which decides everything.

How to run a proper test purchase, five steps.

First: invent the legends. At least three scenarios. The typical case: an ordinary question from an ordinary customer. The atypical one: a situation not covered in the FAQ. And the conflict one: you are unhappy,

you write sharply, you demand a human. It is on the second and third scenarios that bots crumble, and it is exactly these that nobody checks before launch. And in every scenario, check whether the bot issued a receipt and whether the receipt matches what actually happened. A shopper with a receipt in hand writes the report in half an hour instead of half a day: half the evidence is already collected, and by your own bot at that.

Second: prepare a scorecard. What to record: how many steps, how much time, where you stumbled, the bot's verbatim replies, whether it introduced itself as AI, how many times you had to rephrase, whether a human was called and whether the context was passed along. And crucially: next to every item, the corresponding line from the service passport. Remember I said the passport is a ruler? The test purchase is the act of laying the ruler down. Promised "a human within two minutes"? Start the timer.

Third: walk the path for real. Behave like a living person, not a tester: type with typos, start from the middle, ask two questions at once, drop in a voice note. Do not prompt yourself and do not forgive the bot with "well, we haven't finished that part yet." The customer does not know what you have not finished, and does not want to know.

Fourth: record the evidence. Screenshots, screen recording, times by the clock. Not for a court, for the debrief: "the bot was rude" sounds like an opinion, while a screenshot of the rudeness is a fact nobody can argue with at the meeting.

Fifth, the decisive one: consequences. Write this in large letters: a test purchase without follow-up actions is a wasted half hour. Every finding becomes a decision: what we fix, who fixes it, by when. And a month later, a repeat purchase along the same scenario: did it get fixed, or did we just have a nice chat. Purchase, debrief, fix, verify. A cycle, not an event.

Who should do the shopping? In ascending order. The best first shopper is you, the executive: half an hour in the role of your own customer sobers you up better than a quarterly report, as generations of directors can confirm. Next, an employee from another department who does not know "how it is supposed to work." Then an acquaintance who is not in the loop at all. And for regular full checks you need an outside person with a proven method and a fresh eye. You could invite me, the author of this guide. I have, shall we say, some experience.

And finally, the thought this whole section was written for. Every customer of yours is already a mystery shopper. The customer walks your service every day, notices every dead end and every rudeness, keeps a mental scorecard. The only difference is that you never receive the customer's report. Or rather, you do, but in the most expensive form there is: in reviews, in the silence of the departed, and in revenue that "somehow dipped." A test purchase is a way to get the same report earlier, cheaper and without the losses.

Run a test purchase regularly, every six months, and after every major update. The vendor silently updated the model? Your bot may have changed without asking you.

And now the final section of this guide. About the fact that evaluation is not a finish line but a circle: technology changes, I change, and your beautiful rules go stale faster than you think.

5.3. Living Rules: Revise Them Together With Me

Let me open the final section with bad news. Everything you build from this guide will go stale.

Look at who is working against your beautiful documents.

Technology: the vendor updated the model, and the bot started talking differently, asking neither you nor me.

The law: we have seen how requirements arrive in waves, and August 2026 is not the last one.

And finally, me. My expectations grow mercilessly: what delighted me yesterday is the norm today and an irritant tomorrow. A bot answering within a minute felt like a miracle until your competitor started answering in ten seconds. I did not get worse. I simply got used to it. Such is the nature of the most demanding customer, and there is no cure.

The conclusion: AI documents come in exactly two kinds, living and dead. A living document gets revised, it has a last-updated date, and I, by the way, look at that date. A dead one was written, framed and admired. A two-year-old service passport proudly marked version 1.0 tells me one thing: nobody has looked in here since launch. And a bot living by a dead passport lies to me on schedule.

How to keep things alive without bureaucracy. Two mechanisms.

The first: triggers, events after which a revision is mandatory. The vendor updated the model. You launched a new service. A spike of complaints on one topic. The law changed. If any of these happened, the owner of the AI household sits down and checks the passports against reality.

The second: a calendar. Once a quarter, a quick "promises versus facts" check based on test purchases and logs. Once every six months, a major revision: what to remove, what to add, what to change.

Note that the material for revision accumulates by itself: complaints, logs, surveys, test purchases. Speaking of complaints, write down the rule that is worth the most: every justified complaint ends not only with an apology to me but with a change to the policy. An apology treats the case. A policy change treats the cause. A company that apologizes for the same thing for the second year running is apologizing in vain.

And now the main thing. Revise not for me, but with me.

Remember the ladder from Chapter 1: ask, let them try, involve? Back then I was talking about the launch. Well, the ladder does not end at launch. It becomes permanent. A panel of a dozen regular customers who see each new version of the bot first. A beta club with early access. A short conversation with living customers before every major passport revision: what infuriates, what is missing, what should go. Call it whatever you like, a customer council if you want formality. The essence is the same: I take part in editing the rules by which I am served.

What is in it for you, besides my gratitude? This: I am the cheapest and most motivated expert available. A consultant takes the money and leaves. I stay and keep using the service. I notice breakages before your metrics do, because I live inside your service. And I am, damn it, personally invested in everything working: I am the one who has to use it. Give me a legitimate channel of influence, and I will stop expressing my opinions exclusively through reviews.

And so the circle of this guide closes. Appreciate the symmetry. We began with "ask me first." We end with "and never stop asking." Between those two phrases fit everything: the benefit test, the guillotine, the passport, the honest story, the trained team, the living evaluation. AI adoption turned out to be not a project with an end date but a relationship. And relationships, as we know, stay alive exactly as long as the people in them keep talking.

Chapter 5 Checklist. Check Yourself

No.	Question	☐
1	Every metric we collect is tied to a decision: we know what we will do when we see the number	☐
2	We evaluate only what we control	☐
3	The service passport is used as a ruler: every promise is turned into a verifiable indicator	☐
4	Once a quarter customers receive a short survey: was the problem solved, how long did it take, is it simpler than before, what should we fix	☐
5	Logs and the after-dialog micro-rating run continuously, and someone actually reviews them	☐
6	A test purchase is run every six months and after every model update, on three scenarios: typical, atypical, conflict	☐
7	Every test purchase and every justified complaint ends in a change: to the prompt, the passport, the rules	☐
8	All AI documents have a last-updated date, triggers and a revision calendar	☐
9	Customers are involved in revisions: a panel, a beta club or a regular conversation before updates	☐

Nine items. If they are all closed, your AI does not merely work. It grows up together with your customers. That is the rarest quality of all, and unlike models and prompts, competitors cannot copy it.

⁷ Acumen, The Lean Data Field Guide: acumen.org.

Conclusion. Free Up My Chair

Well, that is all, dear executive. Your most demanding customer has had his say.

Let's count what you now have. You know my needs, all three groups, and you know how to ask so that I do not lie. You can count AI's benefits by my yardstick: steps, time, costs, and you have run the guillotine, sending under the blade everything that stole our time for years. Your AI services have acquired passports, your promises have numbers, and your customers have receipts. You talk to me like a human being, you presented the launch as an act of care, and you did not forget to motivate your own employees. And you built an evaluation that looks through my eyes, all the way to test purchases and living rules that grow up together with me.

Not bad for one guide. And notice: almost none of it required a "digital transformation" budget. It required something else, and frankly, something more expensive than money: the willingness to look at your business from my side of the counter.

And memorize the trio like a multiplication table: the policy says how things are done, the passport says what is promised, the receipt says what actually happened. A promise without a passport is vague. A passport without a receipt is unprovable. A receipt without a policy lacks context. Together they turn your AI from a black box into a service you can trust.

Remember Bezos's empty chair from the introduction? Let me reveal the last secret. You do not need to buy that chair. It is already standing in your meeting room. It is just usually covered with folders of vendor presentations. Every time your team argues about models, prompts and budgets, look at that chair and ask the question our whole conversation started with: what will the customer think about all this? Now you know what I would answer. I sit in that chair permanently, even when it looks empty to you.

One confession before we part. I have been harsh with you, snide in places, and I never properly apologized. I will not apologize now either. Because every one of my jabs is actually an invitation: I badly want you to succeed. I am the one who has to use what you build.

There will soon be millions of companies adopting AI. Companies adopting AI in the customer's interest are still a handful. Join the short line. It moves faster.

In the appendix you will find the final checklist: all my requirements from the five chapters, gathered into a single list. Print it, walk through it, count the check marks. Then walk through it again every six months: as you will recall, I grow more demanding by the day.

Thank you for reading to the end and not taking offense. That, by the way, is the best sign that your company is going to be just fine.

Your demanding customer.

P.S. I will be glad to read your feedback: info@aibusiness.vc

Appendix. The Demanding Customer's Final Checklist

All my requirements from the five chapters, gathered into one list. How to use it: print it, walk through it honestly, count the check marks. A check mark is earned only by "done and working," not by "planned for next quarter." Then come back to the list every six months: as you will recall, I grow more demanding by the day.

There are exactly forty-two items. Science fiction fans know: that is the answer to the ultimate question of life, the universe and everything. In our case the question was: "do customers love your AI?"

Chapter 1. Ask the Customer First

No.	Question	☐
1	We know the three groups of customer needs: functional, informational, emotional, and the nine dimensions within them	☐
2	We exported and analyzed complaints and inquiries from the last three months	☐
3	We have the top 10 customer questions, recorded in their own words	☐
4	We looked into the logs: where customers abandon the dialog and after what they leave silently	☐
5	We conducted at least ten interviews by the rules: about the customer's life, about the past, listening	☐
6	We separated facts from compliments, fluff and ideas	☐
7	We chose our step of collaboration: ask, let them try, involve	☐
8	A list of the customer's five main pains, in their own words, hangs in front of the team	☐

Chapter 2. Why Would Your Customer Need Your AI?

No.	Question	☐
9	We counted the customer's three numbers: their steps, their time, their costs. Before adoption and after	☐
10	We honestly identified our outcome: better for the customer, the same, worse	☐
11	We ran the guillotine: removed unnecessary steps and processes before automation, not after	☐
12	We divided what remained between humans and AI, after listening to the people who do the work	☐
13	We measure the customer's time from need to resolution, not the bot's response speed	☐
14	We track repeat contacts and "tell me from the beginning"	☐
15	Someone from management walked the customer's path personally, stopwatch in hand	☐

Chapter 3. Agreeing on the Rules

No.	Question	☐
16	The company has public AI rules: one open document with answers to the ten questions, with a version and a date	☐

17	Every customer-facing AI service has a passport (general provisions, standard, agent procedure), assembled from modules of the rules	☐
18	The passport is published where the customer will find it, written in human language, and fixes: what the bot can and cannot do, deadlines, data, price, a closed list of grounds for refusal	☐
19	The cases when the bot must call a human are spelled out: customer's request, two misunderstandings, money, conflict	☐
20	The bot introduces itself as AI at the start; generated content and deepfakes are labeled (Article 50, deadline August 2, 2026)	☐
21	Employee emotion recognition is switched off; customers are warned about tone analysis	☐
22	Every dialog is logged, and for a significant service the customer automatically receives an interaction receipt with a snapshot of the rules and the passport that were in force	☐
23	The rules and every passport have an owner with a first and last name, and a gap between the document and the bot's behavior is treated as an emergency, not a "nuance"	☐
24	The machine-readable service card is published for customers' AI agents, and it states: a receipt is issued	☐

Chapter 4. Talk to Me Like a Human Being

No.	Question	☐
25	The bot's replies pass the "read it aloud to a neighbor" test: no officialese, the point in the first line, concrete deadlines	☐
26	Customers were told about the AI adoption as an act of care: what it gives them, a human nearby, a named person responsible	☐
27	Clear channels were used: a letter, a post, a notice in the chat, a cheat sheet, an FAQ	☐
28	The company's public AI rules are easy to find from any page of the website, and all communication channels lead to them	☐
29	Changeable information (contacts, names, prices) lives in the public AI rules as a single copy, and the bots pull it from there	☐
30	The team has been told the principle "we cut activities, not employees," and it is backed by deeds	☐
31	Employees are trained: the rules, forbidden data, checking answers, handing a customer to a human (a legal requirement since February 2025)	☐
32	An idea channel is running: what else can we offer customers with AI	☐
33	A head of the AI competence center is appointed: an employee responsible for ideas, bots, passports, training and news	☐

Chapter 5. How to Evaluate AI

No.	Question	☐
34	Every metric we collect is tied to a decision: we know what we will do when we see the number	☐
35	We evaluate only what we control	☐
36	The passport and the receipt are used for double-checking: every promise is turned into a verifiable indicator, and every receipt is a fact to check against	☐
37	Once a quarter customers receive a short survey: was the problem solved, how long did it	☐

	take, is it simpler than before, what should we fix	
38	Logs and the after-dialog micro-rating run continuously, and someone actually reviews them	☐
39	A test purchase is run every six months and after every model update, on three scenarios (typical, atypical, conflict), with a check that the receipt is issued	☐
40	Every test purchase and every justified complaint ends in a change: to the prompt, the passport, the rules	☐
41	All AI documents have a last-updated date, triggers and a revision calendar	☐
42	Customers are involved in revisions: a panel, a beta club or a regular conversation before updates	☐

Count Your Check Marks

0-14. The red zone. Your AI lives its own life, and I live mine, and our relationship rests on my patience. Which, I warn you, is not infinite: I am already eyeing your competitors, and my draft review is ready. The good news: you now know all forty-two steps to fixing it, and the first ones cost zero money.

15-28. The middle. Not bad at all, seriously. It shows that someone thought about me. But I find the holes faster than you patch them, and my loyalty so far rests on habit, not love. Pick the five cheapest of your unclosed items and close them within a month: you will feel the effect before you feel the cost.

29-41. A strong position. I recommend you to friends, and that, as you know, cannot be bought with advertising. Close the remaining items not for the check marks but for the last frontier: they are exactly what separates you from the best.

42 out of 42. A customer champion. I take off my hat and do what a demanding customer almost never does: praise without reservations. I defend your AI in the comments more fiercely than your PR person. One request: do not relax. In six months I will come back with this list. And it may well have grown longer.

How to Work Through the Checklist with Your Team

Do not fill it in alone in your office. Print a copy for everyone at your next team meeting: let each person count the check marks on their own, honestly and without peeking. Then compare. The spread of scores is itself a diagnosis: if the owner counts 30 and the support team counts 12, two different AIs live in your company, mine and yours. Discuss the gaps, pick the five cheapest unclosed items and agree who closes them and by when. Forty-five minutes, and the checklist stops being reading and becomes a plan.

About the Author

Sergei Ponomarev holds a PhD in political science. For more than 20 years he has been making services open, understandable and convenient for the people who use them: first in the public sector, now in business. For 7 years he ran a monitoring program for government and municipal services, conducted hundreds of independent quality assessments and mystery shopping audits, developed an Information Openness Standard for public authorities, completed a fellowship at the Kennan Institute in Washington, DC, and taught at a university for 14 years. For the past year and a half he has been building AI agents hands-on and running the portal aibusiness.vc.

Website: aibusiness.vc

Email: info@aibusiness.vc



Point your camera to open aibusiness.vc